

# Goals

- implement the model  $f_{w,b}$  for linear regression with one variable
- implement and explore the cost function for linear regression with one variable.

## 1. Model Representation

### Notation

Here is a summary of some of the notation you will encounter.

General Notation	Description	Python (if applicable)
$a$	scalar, non bold	
$\mathbf{a}$	vector, bold	
<b>Regression</b>		
$\mathbf{x}$	Training Example feature values	<code>x_train</code>
$\mathbf{y}$	Training Example targets	<code>y_train</code>
$x^{(i)}, y^{(i)}$	$i_{th}$ Training Example	<code>x_i</code> , <code>y_i</code>
$m$	Number of training examples	<code>m</code>
$w$	parameter: weight	<code>w</code>
$b$	parameter: bias	<code>b</code>
$f_{w,b}(x^{(i)})$	The result of the model evaluation at $x^{(i)}$ parameterized by $w, b$ : $f_{w,b}(x^{(i)}) = wx^{(i)} + b$	<code>f_wb</code>

### Tools

In this assignment you will make use of:

- NumPy, a popular library for scientific computing
- Matplotlib, a popular library for plotting data

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
```

## Problem Statement

As in the lecture, you will use the motivating example of diabetes progression prediction. This assignment will use a simple data set with only two data points shown below. These two points will constitute our *data or training set*.

BMI	Diabetes progression
32.1	151
21.6	75

You would like to fit a linear regression model through these two points, so you can then predict diabetes progression for other patients - say, a patient with BMI = 30.5.

Please run the following code cell to create your `x_train` and `y_train` variables. The data is stored in one-dimensional NumPy arrays.

```
In [2]: # x_train is the input variable (BMI)
# y_train is the target (diabetes progression level)
x_train = np.array([32.1, 21.6])
y_train = np.array([151, 75])
print(f"x_train = {x_train}")
print(f"y_train = {y_train}")

x_train = [32.1 21.6]
y_train = [151 75]
```

**Note:** The course will frequently utilize the python 'f-string' output formatting described [here](#) when printing. The content between the curly braces is evaluated when producing the output.

## Number of training examples `m`

You will use `m` to denote the number of training examples. Numpy arrays have a `.shape` parameter. `x_train.shape` returns a python tuple with an entry for each dimension.

`x_train.shape[0]` is the length of the array and number of examples as shown below.

```
In [3]: # m is the number of training examples
print(f"x_train.shape: {x_train.shape}")
m = x_train.shape[0]
print(f"Number of training examples is: {m}")

x_train.shape: (2,)
Number of training examples is: 2
```

One can also use the Python `len()` function as shown below.

```
In [4]: # m is the number of training examples
m = len(x_train)

print(f"Number of training examples is: {m}")

Number of training examples is: 2
```

## Training example $x_i, y_i$

You will use  $(x^{(i)}, y^{(i)})$  to denote the  $i^{th}$  training example. Since Python is zero indexed,  $(x^{(0)}, y^{(0)})$  is (32.1, 151) and  $(x^{(1)}, y^{(1)})$  is (21.6, 75).

To access a value in a Numpy array, one indexes the array with the desired offset. For example the syntax to access location zero of `x_train` is `x_train[0]`. Finish the next code block below to get the  $i^{th}$  training example.

```
In [5]: i = 0 # Change this to 1 to see (x^1, y^1)

x_i = x_train[i] # ith feature value
y_i = y_train[i] # ith target value
print(f"(x^{i}), y^{(i)}) = ({x_i}, {y_i})")

(x^(0), y^(0)) = (32.1, 151)
```

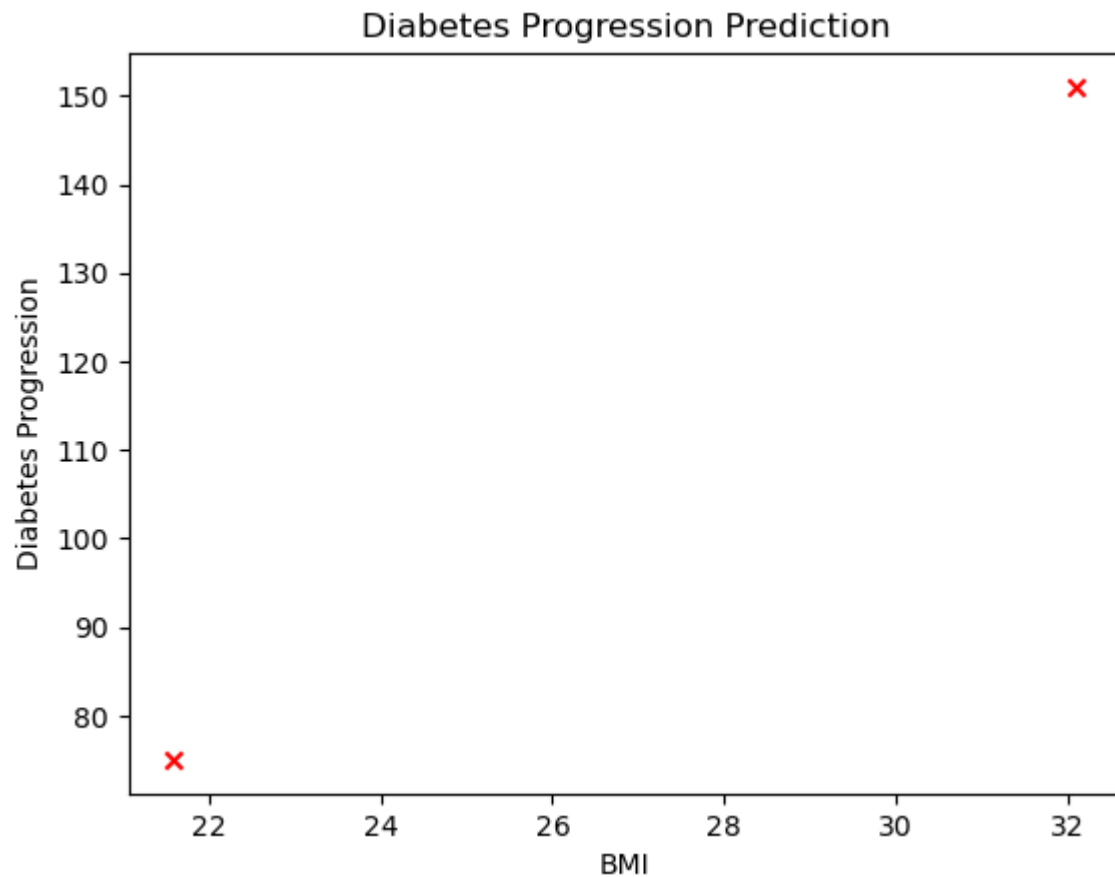
## Plotting the data

You can plot these two points using the `scatter()` function in the `matplotlib` library, as shown in the cell below.

- The function arguments `marker` and `c` show the points as red crosses (the default is blue dots).

You can use other functions in the `matplotlib` library to set the title and labels to display

```
In [6]: # Plot the data points
plt.scatter(x_train, y_train, marker='x', c='r')
# Set the title
plt.title("Diabetes Progression Prediction")
# Set the y-axis label
plt.ylabel('Diabetes Progression')
# Set the x-axis label
plt.xlabel('BMI')
plt.show()
```



## Model function

As described in lecture, the model function for linear regression (which is a function that maps from  $x$  to  $y$ ) is represented as

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b \quad (1)$$

The formula above is how you can represent straight lines - different values of  $w$  and  $b$  give you different straight lines on the plot.

Let's try to get a better intuition for this through the code blocks below. Let's start with  $w = 1$  and  $b = 1$ .

**Note: You can come back to this cell to adjust the model's  $w$  and  $b$  parameters**

```
In [7]: w = 1
b = 1
print(f"w: {w}")
print(f"b: {b}")
```

```
w: 1
b: 1
```

Now, let's compute the value of  $f_{w,b}(x^{(i)})$  for your two data points. You can explicitly write this out for each data point as -

```
for  $x^{(0)}$ , f_wb = w * x[0] + b
```

```
for  $x^{(1)}$ , f_wb = w * x[1] + b
```

For a large number of data points, this can get unwieldy and repetitive. So instead, you can calculate the function output in a `for` loop in the `compute_model_output` function below.

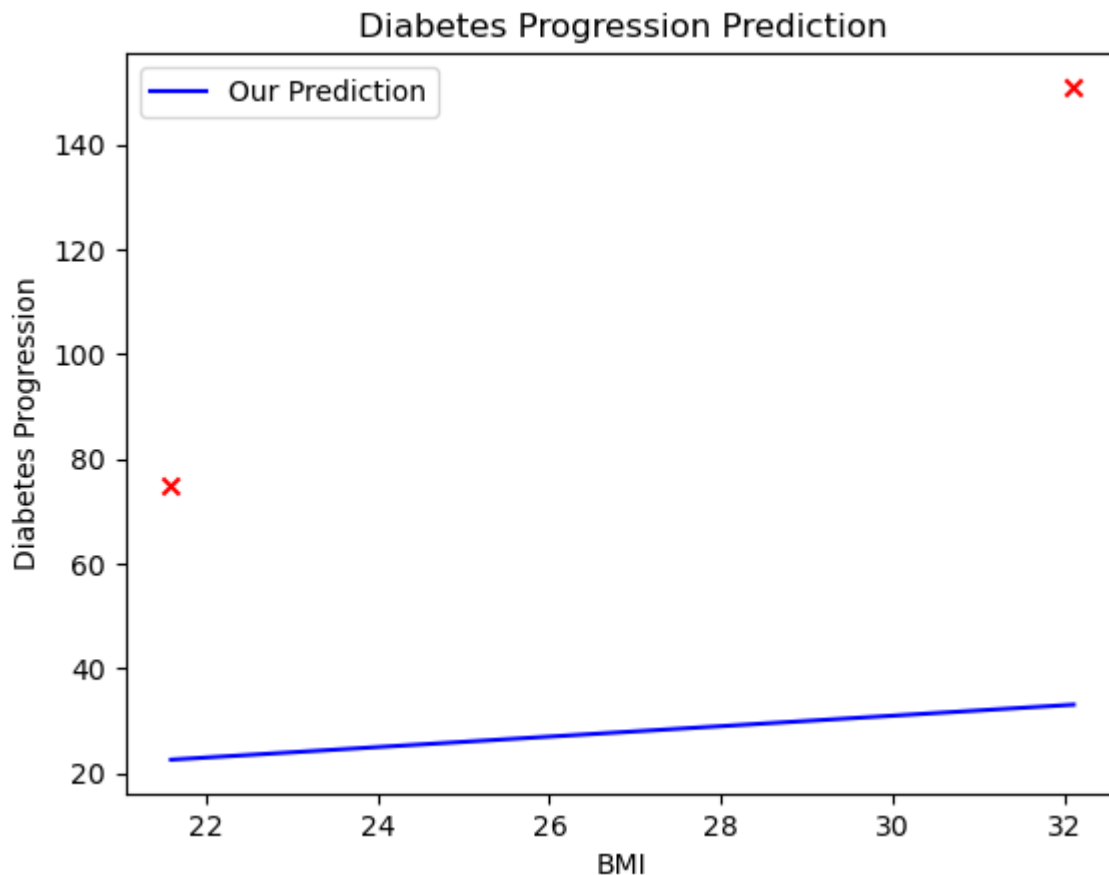
**Note:** The argument description `(ndarray (m,))` describes a Numpy n-dimensional array of shape (m,). `(scalar)` describes an argument without dimensions, just a magnitude.

**Note:** `np.zeros(n)` will return a one-dimensional numpy array with  $n$  entries

```
In [8]: def compute_model_output(x, w, b):  
        """  
        Computes the prediction of a linear model  
        Args:  
            x (ndarray (m,)): Data, m examples  
            w,b (scalar)      : model parameters  
        Returns  
            y (ndarray (m,)): target values  
        """  
  
        m = x.shape[0]  
        f_wb = np.zeros(m)  
        # write a loop to compute f_wb  
        for i in range(0, m):  
            f_wb[i] = w * x[i] + b  
  
        return f_wb
```

Now let's call the `compute_model_output` function and plot the output..

```
In [9]: tmp_f_wb = compute_model_output(x_train, w, b)      # call the compute_model_output fu  
  
        # Plot our model prediction  
        plt.plot(x_train, tmp_f_wb, c='b', label='Our Prediction')  
  
        # Plot the data points  
        plt.scatter(x_train, y_train, c='r', marker='x')  
  
        # Set the title  
        plt.title("Diabetes Progression Prediction")  
  
        # Set the y-axis label  
        plt.ylabel('Diabetes Progression')  
  
        # Set the x-axis label  
        plt.xlabel('BMI')  
  
        plt.legend()  
        plt.show()
```



As you can see, setting  $w = 1$  and  $b = 1$  does *not* result in a line that fits our data.

## Prediction

Try experimenting with different values of  $w$  and  $b$ . What should the values be for a line that fits our data? Note that you can actually compute the theoretical values of  $w$  and  $b$  by hand given the two training examples. Put your best  $w$  and  $b$  in the prediction cell below.

Now that we have a model, we can use it to make our original prediction. Let's predict the diabetes progression of a patient with BMI=30.5. Note: your prediction value should be around 140.

```
In [10]: w = 76/10.5                                # Your best w (7.238)
          b = -32.1*76/10.5 + 151                   # Your best b (-81.343)
          x_i = 30.5
          diabetes_progression = w*x_i + b          # prediction

          print(f"${diabetes_progression:.1f}")
```

\$139.4

## 2. Cost Function

Here, cost is a measure of how well our model is predicting the diabetes progression of a patient.

The equation for cost with one variable is:

$$J(w, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)})^2 \quad (2)$$

where

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b \quad (3)$$

- $f_{w,b}(x^{(i)})$  is our prediction for example  $i$  using parameters  $w, b$ .
- $(f_{w,b}(x^{(i)}) - y^{(i)})^2$  is the squared difference between the target value and the prediction.
- These differences are summed over all the  $m$  examples and divided by  $2m$  to produce the cost,  $J(w, b)$ .

Note, in lecture summation ranges are typically from 1 to  $m$ , while code will be from 0 to  $m-1$ .

The code below calculates cost by looping over each example. In each loop:

- `f_wb`, a prediction is calculated
- the difference between the target and the prediction is calculated and squared.
- this is added to the total cost.

```
In [11]: def compute_cost(x, y, w, b):  
    """  
    Computes the cost function for linear regression.  
  
    Args:  
        x (ndarray (m,)): Data, m examples  
        y (ndarray (m,)): target values  
        w,b (scalar)      : model parameters  
  
    Returns  
        total_cost (float): The cost of using w,b as the parameters for linear regression  
                           to fit the data points in x and y  
    """  
    # number of training examples  
    m = x.shape[0]  
    cost_sum = 0  
  
    # write a loop to compute the summation of the squared difference for all training  
  
    for i in range(0, m):  
        y_i = w*x[i]+b  
        cost_sum += (y_i - y[i])**2  
  
    total_cost = (1 / (2 * m)) * cost_sum  
  
    return total_cost
```

Test your `compute_cost` function using `x_train`, `y_train`, and your best  $w$  and  $b$ . Your total cost should be around 0.

```
In [12]: total_cost = compute_cost(x_train, y_train, w, b) # call compute_cost function
print(f"Total cost is: {total_cost:.1f}")
```

Total cost is: 0.0

Redefine your x\_train and y\_train using a larger training set below. Test your compute\_cost function again.

BMI	Diabetes progression
-----	----------------------

32.1	151
------	-----

21.6	75
------	----

30.5	141
------	-----

22.6	97
------	----

```
In [13]: # Redefining x_train and y_train
x_train = np.array([32.1, 21.6, 30.5, 22.6])
y_train = np.array([151, 75, 141, 97])

total_cost = compute_cost(x_train, y_train, w, b) # call compute_cost function
print(f"Total cost is: {total_cost:.1f}")
```

Total cost is: 27.6

In [ ]: